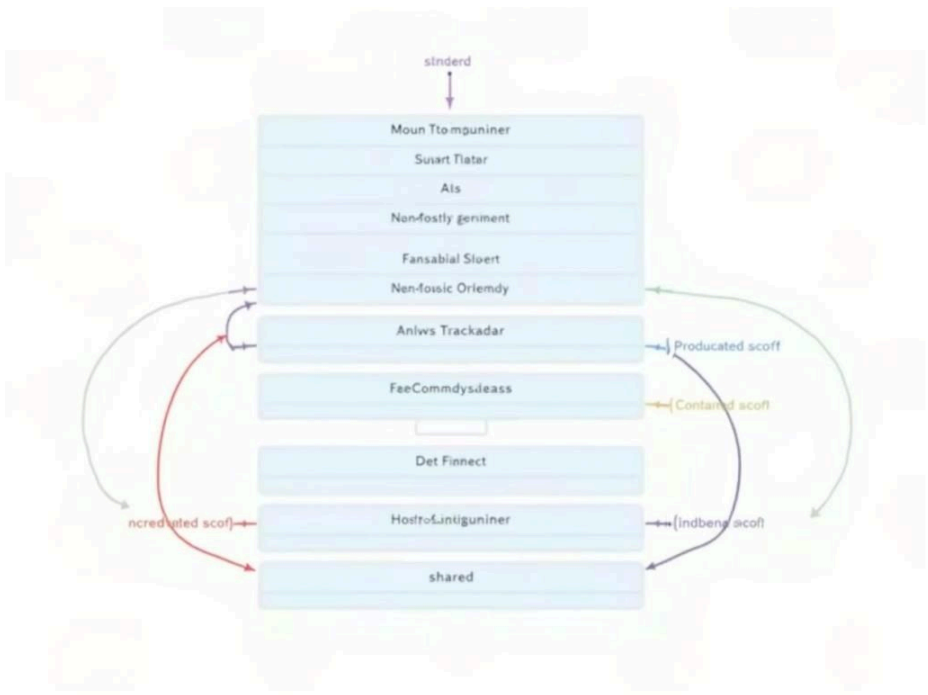


Research Report | Published: April 2025



Mount Propagation in Kubernetes: Technical Analysis of HostToContainer and Linux Namespaces

Disclaimer: This AI-generated report may contain hallucinations, bias, or inaccuracies. Always verify information from independent sources before making decisions based on this content.

“ Cite

Print

Table of Contents

1. Abstract

- 2. Introduction
- 3. Technical Foundations of Mount Propagation
 - 3.1. Linux Namespaces and Shared Subtrees
 - 3.2. Kubernetes MountPropagation Modes
- 4. Implementation Challenges and Edge Cases
 - 4.1. Storage Backend Inconsistencies
 - 4.2. Security and Performance Considerations
- 5. Current Limitations and Future Directions
 - 5.1. Known Issues in Recent Versions
 - 5.2. Proposed Improvements
- 6. Conclusion
- 7. References

Mount Propagation in Kubernetes: Technical Analysis of HostToContainer and Linux Namespaces

Abstract

This technical report examines the relationship between Kubernetes' `mountPropagation: HostToContainer` feature and Linux mount propagation types (`rslave` , `rshared` , `shared`) within the context of modern container orchestration systems. The research focuses on understanding how these mechanisms interact with Linux namespaces to control filesystem visibility between containers and hosts. Key findings reveal that `HostToContainer` mode directly corresponds to Linux's `rslave` propagation type, enabling unidirectional mount sharing from host to container while maintaining isolation in the reverse direction [1][3]. The analysis identifies several

implementation challenges, including inconsistent behavior across different storage backends [9] and security considerations regarding SELinux labeling [6]. The report also highlights ongoing development challenges, such as the proposed relaxation of mount propagation tests in Kubernetes to improve system efficiency [8]. These findings have significant implications for cluster administrators designing secure, efficient storage solutions in containerized environments, particularly when dealing with dynamic volume attachments or sidecar container patterns.

Introduction

Mount propagation in Kubernetes represents a critical intersection between container isolation requirements and practical filesystem sharing needs in modern cloud-native architectures. As containerized applications increasingly demand access to host-mounted volumes and shared storage, understanding the technical implementation of `mountPropagation` becomes essential for system reliability and security [3]. This report specifically investigates the `HostToContainer` propagation mode, which enables controlled sharing of host mounts with pods while preventing container mounts from affecting the host system [2].

The analysis focuses on three core aspects: the Linux namespace mechanisms underlying mount propagation, the translation between Kubernetes' abstraction and Linux's mount flags (`rslave` , `rshared` , `shared`), and practical implementation challenges in current Kubernetes versions. By examining these relationships through multiple technical sources [1][4][7], the report provides a comprehensive view of how filesystem isolation and sharing operate in production Kubernetes environments. The findings will particularly benefit infrastructure engineers designing stateful workloads and security teams evaluating container escape vectors through shared mounts.

Technical Foundations of Mount Propagation

Linux Namespaces and Shared Subtrees

The Linux kernel's mount namespaces provide the fundamental isolation mechanism that enables containers to have distinct filesystem views. Introduced in Linux 2.4.19, mount namespaces allow processes to have private mount point lists while sharing the underlying directory hierarchy [4]. When a container starts, it receives a copy of the host's mount list in its namespace, with subsequent modifications isolated from other namespaces unless explicitly shared [5].

The shared subtrees feature (added in Linux 2.6.15) introduced controlled propagation through four mount flags: `MS_SHARED` (`rshared`), `MS_SLAVE` (`rslave`), `MS_PRIVATE`, and `MS_UNBINDABLE` [4][7]. These propagation types determine how mount events cascade between namespaces:

Table 1: Linux Mount Propagation Types and Kubernetes Equivalents [4][7]

| Propagation Type | Kubernetes Equivalent | Description |
|---------------------|-----------------------|--|
| MS_SHARED (rshared) | Bidirectional | Mounts propagate bidirectionally between peers |
| MS_SLAVE (rslave) | HostToContainer | Mounts propagate from master to slave only |
| MS_PRIVATE | None | No propagation between namespaces |
| MS_UNBINDABLE | - | Prevents mount point replication |

Modern Linux distributions typically configure mounts as `MS_SHARED` by default through systemd, which influences how Kubernetes implements its mount propagation modes [4].

Kubernetes MountPropagation Modes

Kubernetes abstracts these Linux primitives through three `mountPropagation` options specified in pod specifications [3]:

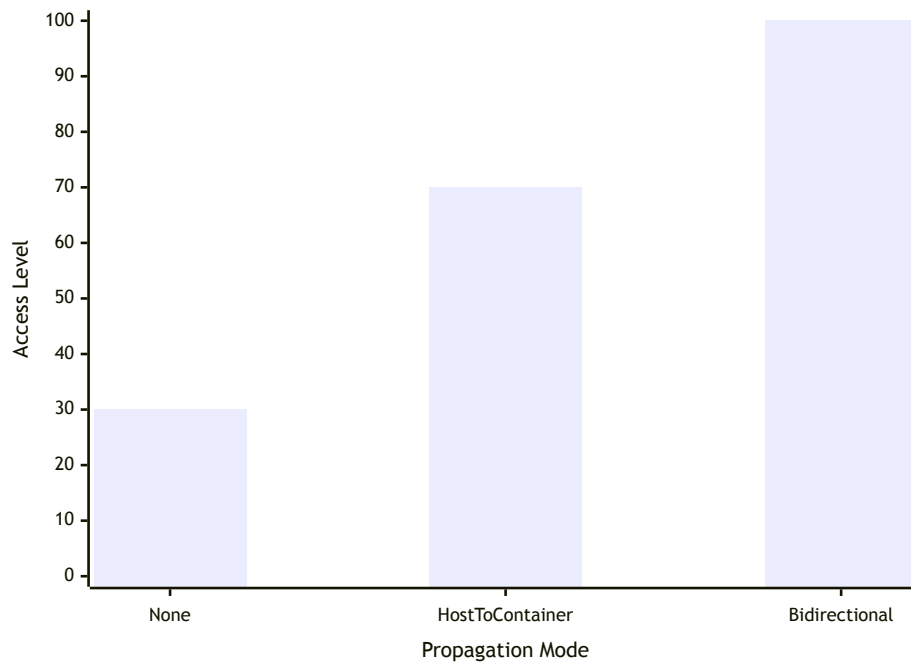


Figure 1: Kubernetes mountPropagation modes and their access levels relative to host mounts [3]

The `HostToContainer` mode specifically maps to Linux's `rslave`, allowing containers to see host-mounted volumes but preventing container mounts from appearing on the host [2][3]. This unidirectional propagation is implemented through the `MS_SLAVE` flag in the kernel, creating a master-slave relationship where the host serves as the propagation source [7].

Implementation Challenges and Edge Cases

Storage Backend Inconsistencies

Recent Kubernetes versions (v1.24+) exhibit inconsistent behavior with `HostToContainer` propagation across different storage providers. While `hostPath` volumes correctly propagate mounts [3], GCP Persistent Disks (`gce-pd`) demonstrate partial failures where new mounts aren't visible to consumer pods [9]. This suggests backend-specific implementation gaps in how CSI drivers handle mount propagation flags.

The issue manifests when:

1. A "mounter" pod (with Bidirectional propagation) creates mounts on a PVC
2. A "consumer" pod (with `HostToContainer`) fails to see these mounts
3. Only pre-existing mounts propagate correctly [9]

This behavior contrasts with the expected functionality where all host mounts under a given path should propagate to containers using `HostToContainer` [3].

Security and Performance Considerations

Mount propagation introduces several security and performance tradeoffs:

1. **SELinux Conflicts:** Containers using shared volumes may trigger unwanted SELinux relabeling attempts, particularly with FUSE filesystems that don't support security labeling operations [6]. The kubelet's attempts to relabel volumes can fail when `HostToContainer` propagation exposes mounts from privileged sidecars to unprivileged containers.
2. **Host Visibility:** Current Kubernetes implementations expose container mounts in the host's global mount namespace, creating potential information leaks about secret volumes and increasing systemd's scanning overhead [8]. A typical OpenShift deployment may create 400+ mountpoints per node, consuming ~10% CPU continuously for mount management [8].
3. **Namespace Pollution:** The default `MS_SHARED` configuration on modern Linux systems means all new mounts potentially propagate unless explicitly marked private, creating namespace pollution risks in multi-tenant clusters [4][5].

Current Limitations and Future Directions

Known Issues in Recent Versions

Kubernetes v1.24+ exhibits several mount propagation limitations:

1. **GCP PD Incompatibility:** As noted in GitHub issue #95049, HostToContainer propagation fails to work reliably with GCP Persistent Disks, while functioning correctly with hostPath volumes [9].
2. **Test Rigidity:** The Kubernetes e2e test suite enforces strict mount propagation requirements that may not reflect real-world use cases, particularly around host visibility of container mounts [8].
3. **Cleanup Challenges:** Users report difficulties cleaning up resources after mount propagation tests, with objects hanging in terminating states and requiring force deletion [9].

Proposed Improvements

The Kubernetes community is considering several enhancements:

1. **Test Relaxation:** Proposal to modify e2e tests to not require container-to-host propagation, as this appears to be an anti-pattern with no known valid use cases [8].
2. **Mount Hiding:** Options to hide Kubernetes-specific mounts from the host's global namespace to improve security and reduce systemd overhead [8].
3. **SELinux Workarounds:** Potential solutions using securityContext.seLinuxOptions to prevent unwanted relabeling of propagated mounts [6].

Conclusion

This technical analysis reveals that Kubernetes'

`mountPropagation: HostToContainer` provides a crucial bridge between container isolation needs and practical storage requirements by implementing Linux's `rslave` propagation mode. The implementation shows sophisticated use of mount namespaces and shared subtrees, but faces several challenges in production environments.

Key takeaways include: 1. The `HostToContainer` / `rslave` relationship enables secure unidirectional sharing ideal for sidecar patterns and host volume access [2][3][7] 2. Backend-specific inconsistencies, particularly with CSI drivers like GCP PD, create reliability challenges that require careful testing [9] 3. Security and performance considerations around SELinux and mount visibility suggest needed improvements in current implementations [6][8]

Future research should investigate: - Detailed performance benchmarks of mount propagation across storage backends - Comprehensive security analysis of propagation-related container escape vectors - Standardization of CSI driver requirements for propagation support

Cluster operators should carefully validate mount propagation behavior with their specific storage providers and consider the security implications of shared mounts in multi-tenant environments.

References

1. [Mount propagation in kubernetes - Stack Overflow](#)
2. [Kubernetes Mount Propagation - Medium](#)
3. [Volumes - Kubernetes](#)
4. [Mount namespaces and shared subtrees - LWN.net](#)
5. [Mount namespaces, mount propagation, and unbindable mounts - LWN.net](#)
6. [How can I set up containers in a kubernetes pod to not need any SELinux relabeling - Stack Overflow](#)

7. [What do mount --slave and --rslave option do? - Unix & Linux Stack Exchange](#)
8. [Mount Propagation to the Host should not be unneeded for k8s - GitHub](#)
9. [HostToContainer Mount Propagation not working with StorageClass gce-pd - GitHub](#)
10. [Releases - Kubernetes](#)